LINK ANALYSIS

Determining Author Popularity with Page Rank Algorithm using Apache Spark

Data Science For Economics

Aleksandra Zografska 48924A

"I/We declare that this material, which I/We now submit for assessment, is entirely my/our own work and has not been taken from the work of others, save and to the extent that such work has been cited and acknowledged within the text of my/our work, and including any code produced using generative AI systems. I/We understand that plagiarism, collusion, and copying are grave and serious offences in the university and accept the penalties that would be imposed should I engage in plagiarism, collusion or copying. This assignment, or any part of it, has not been previously submitted by me/us or any other person for assessment on this or any other course of study."

Dataset

For this link analysis project, I have chosen the arXiv dataset which contains a snapshot of around 1.7 million academic articles that are frequently updated. I obtained it by accessing Kaggle with a personal API key, at the following link. The dataset is distributed as a JSONL file named arxiv-metadata-oai-snapshot.json. The full dataset at the moment of running my project is approximately 4GB. The full access to the code can be found on GitHub: https://github.com/Zografska/page-rank-arxiv.

As the dataset is really big, it took a lot of computing power to load it initially into memory. To avoid this problem, I load the dataset as a Resilient Distributed Dataset (RDD) into Spark, and set the number of the lower bound of partitions to 8. Since spark is evaluating lazily, the file is only read when an action like .count(), .collect(), or .take() is triggered. Because the other calculations were taking a lot of time, I decided to introduce a variable that samples a percentage of the json lines from my dataset. To ensure consistency in the result, I also set a seed for the sampling. The RDD object has a direct sampling function that I utilized, where the sample portion is controlled through a variable. To make the results more meaningful and focused, I decided to add some filtering based on the paper category.

Data preprocessing

In order to do the link analysis, first I needed to obtain the author names. Name extraction and indexing was done in a first pass over the RDD using a method that collected a list of author names for every paper and then flatMap, distinct, and zipWithIndex, which assigned a unique integer index to every distinct author name. The name_to_idx dictionary is sent from the driver to every worker node once, so each worker can resolve author names to integer IDs with a local in-memory lookup. The alternative, joining the pairs RDD against a distributed lookup table, would require shuffling data across the network multiple times, once per join, which is far more expensive at scale.

In the second pass, I extracted the authors for each paper, then all combinations of co-authors were generated using Python's itertools.combinations, and both directions of each pair were added to represent the co-authorship graph as undirected to ease calculation of the connection matrix down the line. The resulting pairs RDD is then flat-mapped, deduplicated with distinct, repartitioned to even out the workload of the workers, and cached using Spark function chaining.

In this chaining I also set the number of partitions proportional to the number of pages I have loaded up by calculating the maximum of the number of cpu cores times 4 and the number of papers over 50,000 which in turn would give 1 partition to every 50,000 papers. I obtained this number by consulting the Spark tuning guidelines, which suggest one worker node for every 100-200KB of data. Taking that the average size of a parsed data object is between 2 to 4 KB, I took 50,000 as my number. This ensures the partition count scales with dataset size, preventing both over-partitioning on small datasets and bottlenecks on large ones. In the case of the free Tier in GoogleColab, there are only 2 virtual CPUs meaning we can run the operations at 2 partitions at a time. I discovered the 8 partition minimum threshold, by empiristically running the algorithm with a different number of partitions. This empirical evidence coincides with the general Spark recommendation for partitions that's 2-4 partitions per core, in order to improve Spark performance because the scheduler can pipeline tasks: if one task stalls due to garbage collection or skewed computation, another partition can immediately run on the same core instead of leaving it idle.

Algorithms and processing

I implemented the PageRank algorithm using the power iteration method. The co-authors graph is represented as a connection matrix, where each entry holds a source node B, a target node A, and the corresponding transition weight equal to 1 divided by the number of possible transitions from source node B. This matrix is parallelized into a Spark RDD and cached in memory, since it is reused on every iteration.

Before the iterations starts, I initialize a list of 1/n values, with length corresponding with the number of nodes, which is going to be the starting PageRank

of the nodes. Each iteration starts with the current PageRank vector broadcast to all Spark workers. For each edge (target, source, weight) in the RDD, the contribution of the source node to the target node is computed as weight * PageRank[source]. These contributions are then summed per target node using reduceByKey, and the results are collected back to the new page rank variable as a dictionary. The broadcast variable is unpersisted after each iteration to free worker memory.

The new PageRank vector is then assembled from this dictionary and teleportation with the damping factor is applied:

$$PR(i) = d * \sum_{j \to i} \frac{PR(j)}{out(j)} + \frac{(1-d)}{n}$$

where d=0.85 is the damping factor and n is the total number of authors. The (1-d)/n represents the random surfer jumping to any node at random, which ensures convergence and handles spider traps i.e the surfer being stuck in a circle of authors that only collaborate between each other.

Convergence is checked at each iteration by computing the L1 distance between the old and new PageRank vectors. Iteration stops when this distance falls below a tolerance of 10ˆ6 or after a maximum of 100 iterations, both of which are arguments that could be configured in the function call.

Experiments and Results

Initially, I did an effort to find a way to process the whole dataset. Even though I successfully found a way to read the data in chunks, the resource and time constraints were still a problem since it called for an estimated 2 hour processing just for creating the author links. I estimated this number by looking into the active job in the Spark UI, which I had to tunnel with ngrok to make it accessible to my browser.

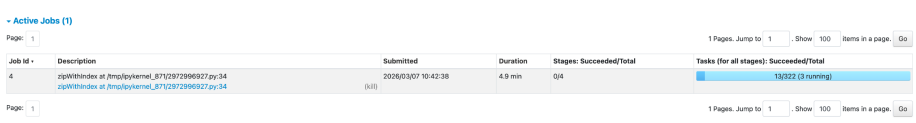

To get around the physical and time constraints, I decided to sample the data, which is handled in my case by spark itself. The variable SUBSAMPLE_FRAC can be adjusted to any value in the notebook and then with the sample() method call on the RDD object, the sampling will be performed when an action is called. To ensure sampling consistency, I pass a seed parameter to this method call.

When a fraction f of nodes is retained, edge survival probability drops to f^2, since an edge only survives if both endpoints are sampled. Keeping 10% of nodes therefore retains roughly 1% of edges. This means prolific authors may appear with only a handful of connections in the sample, underestimating their

PageRank, and dense collaboration clusters can fragment when bridge authors are not selected. Nonetheless, my assumption is that the thight-knit collaboration clusers can survive this sampling, preserving the general structure of collaboration communities.

Given this disclaimer, I would like to discuss a few different results I got from running the algorithm with different values:

Run1: No filter, sampling of 1% of the data The 1% sample yielded 29,542 papers, from which 104,221 unique authors and 9,489,701 co-authorship links were extracted. The algorithm converged after 55 iterations.

Rank	Name	Score	Rank	Name	Score	Rank	Name	Score
1	Liu Yang	0.000161	8	Chen Xi	0.000082	15	Liu Wei	0.000076
2	Li Y.	0.000101	9	Wang Wei	0.000082	16	Wang Z.	0.000074
3	Chen Wei	0.000101	10	Liu Feng	0.000081	17	Chen Bin	0.000074
4	Li Wei	0.000097	11	Zhang Wei	0.000080	18	Wang Lei	0.000073
5	Watanabe Kenji	0.000096	12	Zhang Yang	0.000077	19	Wang Xin	0.000073
6	Taniguchi Takashi	0.000093	13	Chen Hao	0.000077	20	Zhang Bo	0.000073
7	Zhang Y.	0.000086	14	Wang Y.	0.000077			

The top-ranked author is Liu Yang with a PageRank score of 0.000161, noticeably ahead of the next authors, which clusters tightly between 0.000082 and 0.000101. The remainder of the top 20 is dominated by common Chinese names such as Li Y., Chen Wei, Zhang Y., and Wang Wei. These are highly ambiguous name strings that aggregate publications from many distinct individuals who were never disambiguated in the source data. A single "Chen Wei" node in the graph therefore accumulates co-authorship edges from potentially dozens of unrelated researchers, artificially inflating its PageRank. This assumption is further confirmed by searching this name on Google Scholar which resulted with a search result that's 10 pages long.

That said, the results from this rank cannot be strictly taken, and a claim like "Liu Yang is the most active in the research community" **cannot** be concretely supported.

Here the PageRank scores are inherently small because they form a probability distribution summing to 1 across all 104,221 nodes, so the meaningful metric is not the absolute value but how many times larger a score is than the uniform baseline of around 0.0000096.

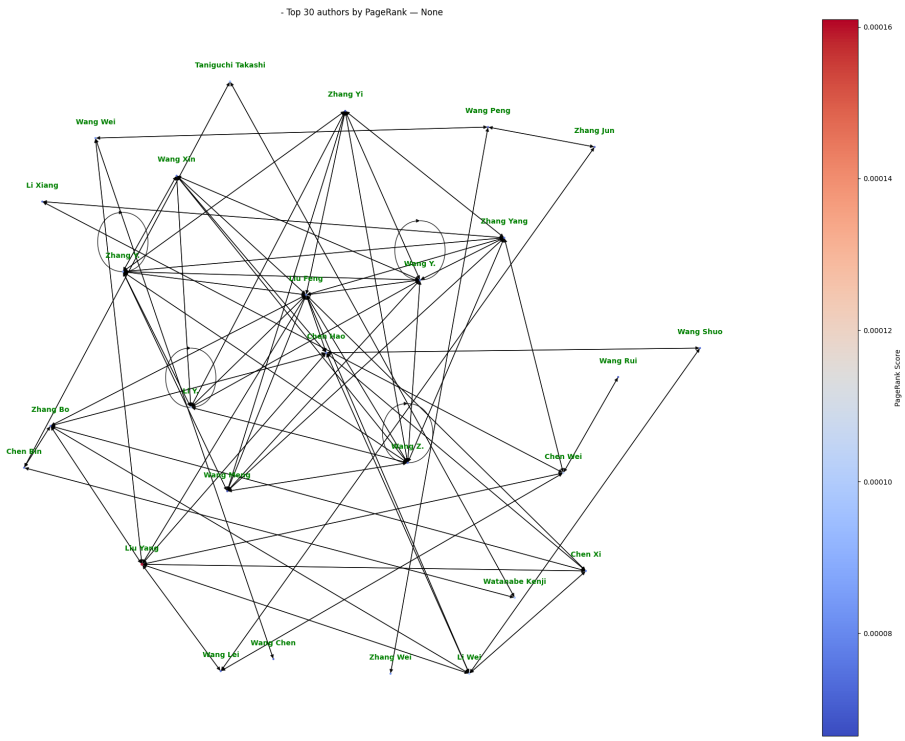

Run2: CS filter, sampling of 1% of the data Filtering to only computer science papers and on the same 1% sample converged in 49 iterations. The first thing to note is that the scores are roughly twice as large as in Run 1: the top author Yang Liu scores 0.000336 compared to 0.000161 previously. This is expected because the graph is smaller after filtering, so the fixed probability mass of 1 is distributed across fewer nodes, pushing individual scores up.

The results are also more interpretable, since names like Philip S. Yu and Dacheng Tao appear, who are both famous CS and machine learning researchers. Their presence near the top suggests the CS filter is doing meaningful work in surfacing real influence rather than name collision results. That said, names like Wei Wang, Wei Zhang, and Hao Wang are still present enough to carry some ambiguity, and appear as the result of the before mentioned name collision.

Rank	Name	PageRank Score	Rank	Name	PageRank Score	Rank	Name	PageRank Score
	Yang			Min			Yu	
1	Liu Lei	0.000336	8	Zhang	0.000202	15	Wang Yu	0.000185
2	Zhang Hao	0.000296	9	Xiang Li Rui	0.000196	16	Zhang Wei	0.000181
3	Wang Wei	0.000230	10	Wang Dacheng	0.000193	17	Wei Philip	0.000172
4	Chen Wei	0.000225	11	Tao	0.000193	18	S. Yu Xinyu	0.000170
5	Zhang Wei	0.000209	12	Bo Li	0.000191	19	Zhang Yang	0.000162
6	Wang Jie	0.000208	13	Yi Wang Yong	0.000191	20	Wang	0.000160
7	Zhang	0.000205	14	Liu	0.000186			

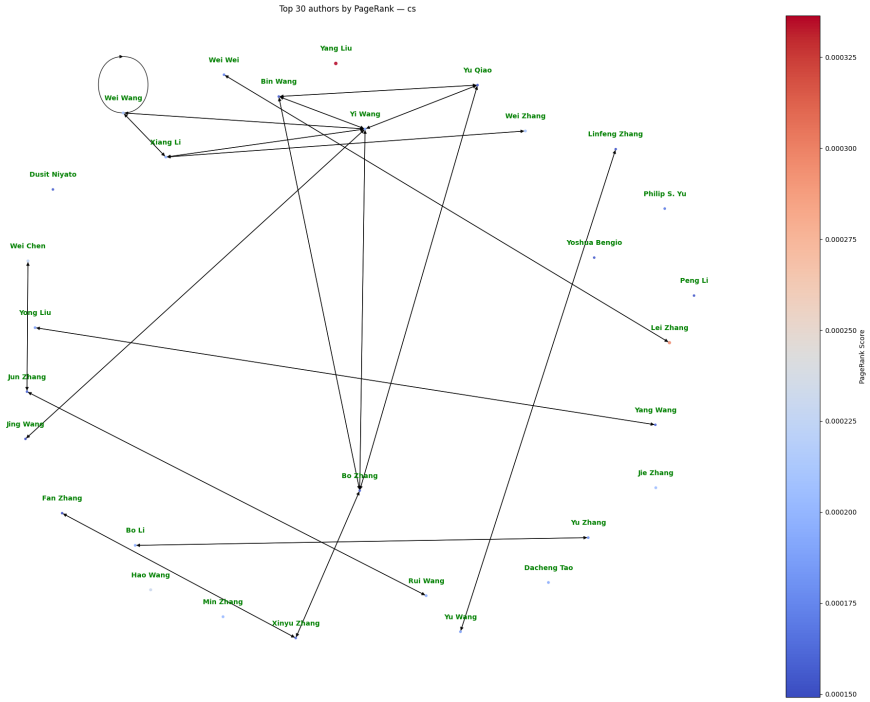

The graph tells a clear story about sampling sparsity. Most of the 30 nodes appear as isolated dots with no visible edges at all, meaning their co-authors were not retained in the 1% sample. This illustrates a sampling fragmentation problem: even at the top of the PageRank ranking, the sampled subgraph captures only a thin slice of the true collaboration structure.

Let's see what happens if we sample 10% of the data applying the same filter.

Run3: CS filter, sampling of 10% of the data Moving from 1% to 10% sampling scales the graph to 207,345 nodes and 3,005,283 edges which is roughly 10x more nodes but 30x more edges. Since the sample size is larger, there is a bigger chance for an author relationship to survive the sampling, therefore edges recover quickly. The algorithm converged in 53 iterations, in line with previous runs.

Top PageRank scores drop slightly (from 0.000336 to 0.000296) because the larger node count. Liu Yang reappears at the top and the ranking is again dominated by ambiguous Chinese name strings, with recognizable figures like Dacheng Tao and Philip S. Yu dropping out. The 10% CS sample is pulling in a much broader slice of the dataset, likely reintroducing the name-collision problem that the tighter 1% sample had partially suppressed.

Rank	**Name**	**Score**	**Rank**	**Name**	**Score**	**Rank**	**Name**	**Score**
1	Liu Yang	0.000296	8	Zhang Lei	0.000155	15	Liu Wei	0.000134
2	Wang Wei	0.000221	9	Wang Jun	0.000154	16	Liu Chang	0.000130
3	Zhang Yu	0.000180	10	Wang Hao	0.000153	17	Zhang Hao	0.000125
4	Li Xiang	0.000178	11	Li Wei	0.000146	18	Li Yang	0.000124
5	Chen Hao	0.000167	12	Wang Yu	0.000143	19	Wang Yan	0.000123
6	Wang Xin	0.000158	13	Zhang Wei	0.000141	20	Yang Fan	0.000123
7	Liu Xin	0.000158	14	Li Bo	0.000141			

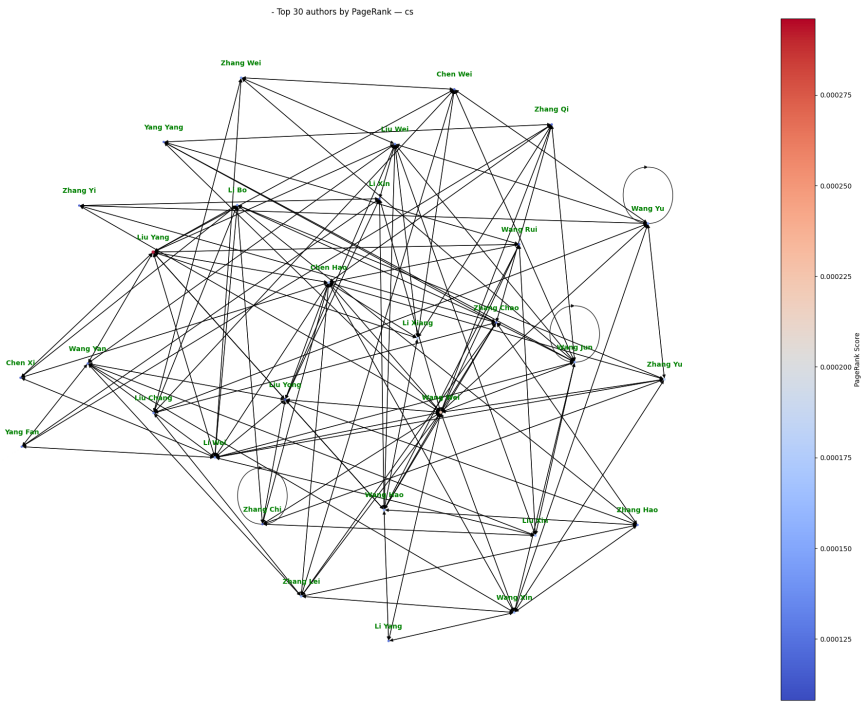

I wanted to change the filter to something that had less published papers, to see how the algorithm would work there.

Run4: Econ filter, sampling of 1% of the data Filtering to economics at 1% gives a tiny graph of only 320 authors and 672 links, converging in 16 iterations. Economics is just much less represented in the dataset than CS, so the small size is expected.

Rank	Name	Score	Rank	Name	Score	Rank	Name	Score
1	Yuan Liao	0.005012	8	Michael Kummer	0.003258	15	Aidos Joldassov	0.003125
2	Stephen Watson	0.004438	9	Marit Hinnosaar	0.003258	16	Patrick Stokkink	0.003125
3	Toomas Hinnosaar	0.004438	10	Bakytzhan Amralinova	0.003125	17	Andres Santos	0.003125
4	José Carlos R. Alcantud	0.003258	11	Aliya Turegeldinova	0.003125	18	Sheng Xu	0.003125
5	Alfio Mate Giarlotta	0.003258	12	Miklos Fodor	0.003125	19	Jonathan Berrisch	0.003125
6	Domenico Cantone	0.003258	13	Akerkin Eraliyeva	0.003125	20	Shao-Yu Jheng	0.003125
7	Olga Slivko	0.003258	14	Chen Dayou	0.003125			

The scores are an order of magnitude higher than previous runs, with the top author scoring 0.005012. This is a direct consequence of the graph size: with 320 nodes the uniform baseline is $1/320 = 0.003125$, and I can

see that everything from rank 4 downward sits exactly at that value. What I find most interesting here is the name quality. Unlike the CS runs, the top authors are unambiguous individual researchers Marit Hinnosaar, Toomas Hinnosaar, Michael Kummer, Stephen Watson. This is the first run that doesn't have name collision noise, which makes the econ rankings the most trustworthy of all runs in terms of actually reflecting real people.

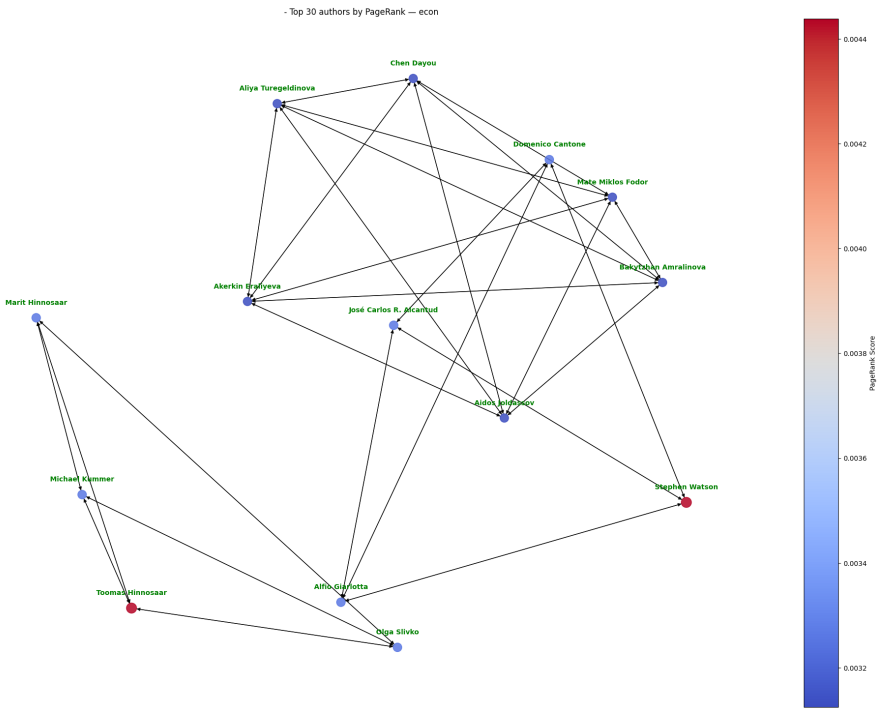

I can also see a visible cluster in the upper portion of the graph connecting Chen Dayou, Aliya Turegeldinova, Domenico Cantone and others, which likely represents a real co-authorship group that survived sampling intact.

Run5: Econ filter, sampling of 10% of the data At 10% sampling I get 3,186 authors and 9,040 links. The score distribution is also healthier than Run 4 because instead of most nodes sitting at the uniform baseline, I see a meaningful gradient from Chernozhukov at 0.001673 downward, suggesting real structural differentiation.

Rank	Name	Score	Rank	Name	Score	Rank	Name	Score
	Chernozhukov			Huber			Bai	
1	Victor	0.001673	8	Florian	0.000754	15	Jushan	0.000694
	Syrgkanis			Liao			Wager	
2	Vasilis	0.000947	9	Yuan	0.000749	16	Stefan	0.000687
	Pin			Li			Ugander	
3	Paolo	0.000932	10	Degui	0.000746	17	Johan	0.000685
							Wang	
	Athey			Ortega			Wein-	
4	Susan	0.000902	11	Josue	0.000739	18	ing	0.000681
	Imbens			Wang				
5	Guido	0.000833	12	Chao	0.000705	19	Ball Ian	0.000679
				Strittmatter				
	Zaccaria			An-			Ma	
6	Andrea	0.000764	13	thony	0.000702	20	Xinwei	0.000659
	Lee			Tamuz				
7	Sokbae	0.000755	14	Omer	0.000695			

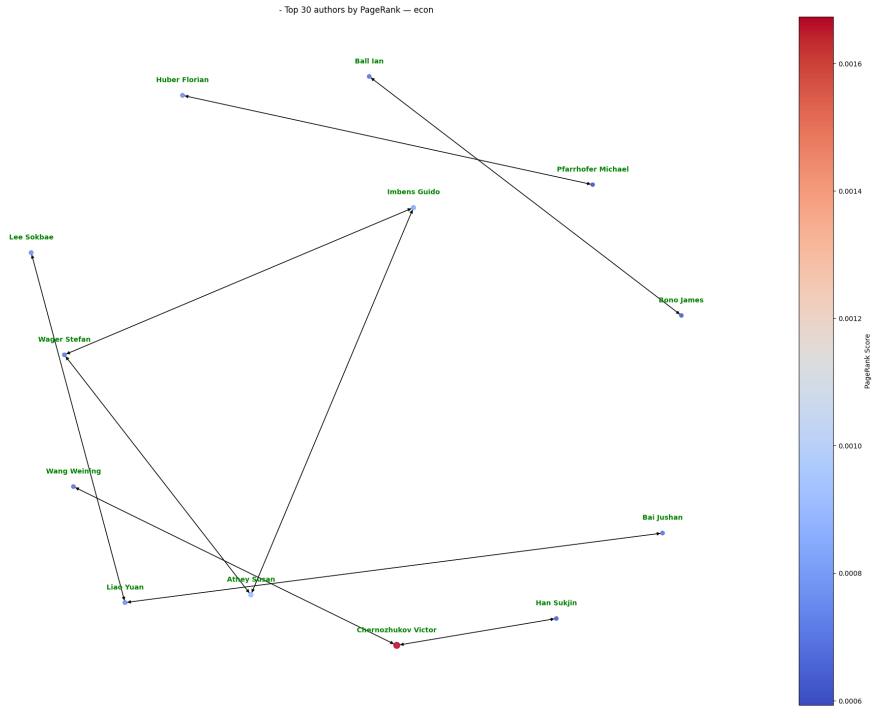

The graph still shows fragmentation. The top-30 visualization is sparse, with most nodes weakly connected and only a loose central cluster around Athey, Imbens, Chernozhukov and Wager, which is a well-known real collaboration group in the literature. Convergence took 59 iterations, the most of any run,

which I attribute to the larger and more structurally complex graph requiring more steps to stabilize.

Conclusion

Across all five runs, name collision was the main issue, not sampling rate. The CS runs were mostly uninterpretable at both scales, while the econ runs surfaced real recognizable researchers even at 1%. When the data is clean, PageRank works: the cluster in Run 5 is exactly what I would expect. The biggest takeaway is domain filtering and name disambiguation matter more than scale i.e. a small, clean graph produces more meaningful rankings than a large, noisy one.